

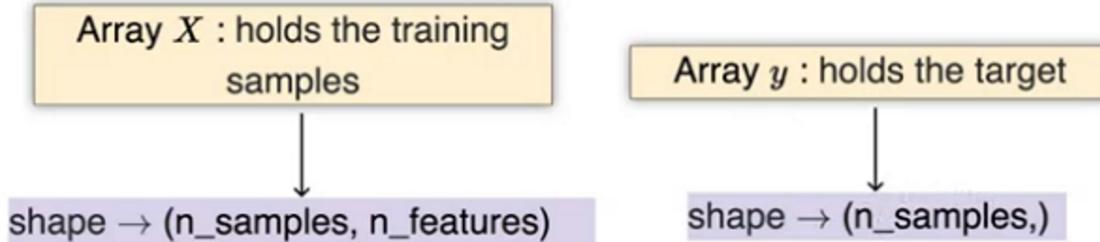
Neural Network

27 March 2023 10:41 AM

MULTILAYER PERCEPTRON (MLP)

- It is a supervised learning algorithm.
- MLP learns a **non-linear function approximator** for either classification or regression depending on the given dataset.
- In `sklearn`, we implement MLP using:
 1. `MLPClassifier` for classification
 2. `MLPRegressor` for regression
- `MLPClassifier` supports **multi-class classification** by applying Softmax as the output function.
- It also supports **multi-label classification** in which a sample can belong to more than one class.
- `MLPRegressor` also supports **multi-output regression**, in which a sample can have more than one target.

Training data





- How to implement MLPClassifier?

Step 1: Instantiate a **MLP** classifier estimator.

```
1 from sklearn.neural_network import MLPClassifier COPY  
2 MLP_clf = MLPClassifier()
```

Step 2: Call **fit** method on **MLP** classifier object with **training feature matrix** and **label vector** as arguments.

```
1 # Model training with feature matrix X_train and  
2 # label vector or matrix y_train  
3 MLP_clf.fit(X_train, y_train)
```

MLPClassifier



Step 3: After fitting (training), the model can make predictions for new samples (**X_test**) using two methods:

```
1 MLP_clf.predict(X_test)  
2 MLP_clf.predict_proba(X_test)
```

predict →

- gives labels for new samples
- for example:
`array([1, 0])`

predict_proba →

- gives vector of probability estimates per sample
- for example:
`array([1.967...e-04, 9.998...-01])`
- MLPClassifier supports only the **Cross-Entropy loss function**

hidden_layer_sizes

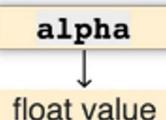
- This parameter sets the number of layers and the number of neurons in each layer.
- It is a **tuple** where **each element in the tuple represents the number of neurons** at the i th position where i is the index of the tuple.
- The **length of tuple** denotes the **total number of hidden layers** in the network.

To create a 3 hidden layer neural network with 15 neurons in first layer, 10 neurons in second layer and 5 neurons in third layer:

```
1 MLPClassifier(hidden_layer_sizes=(15,10,5))COPY
```

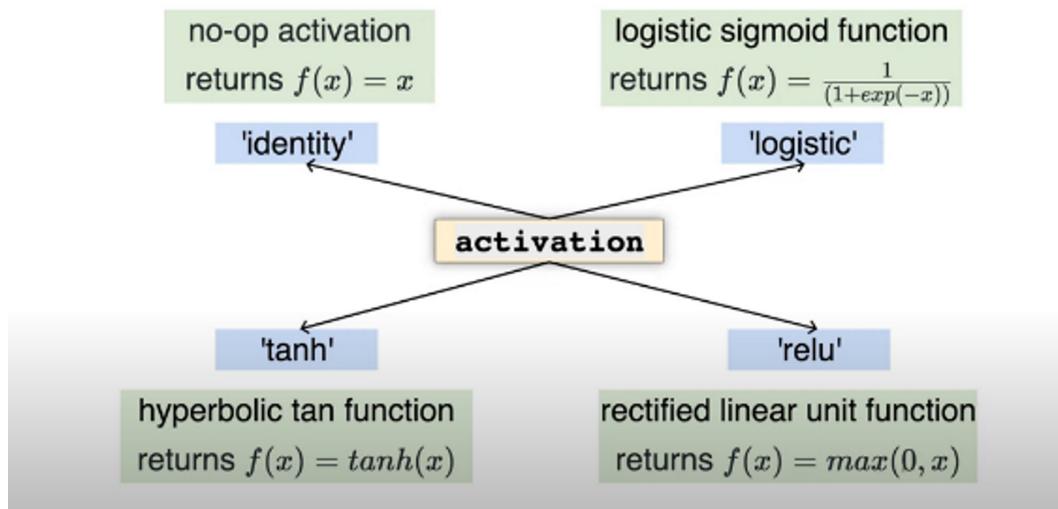
How to perform regularization in MLPClassifier?

- The alpha parameter sets L2 penalty Regularization parameter

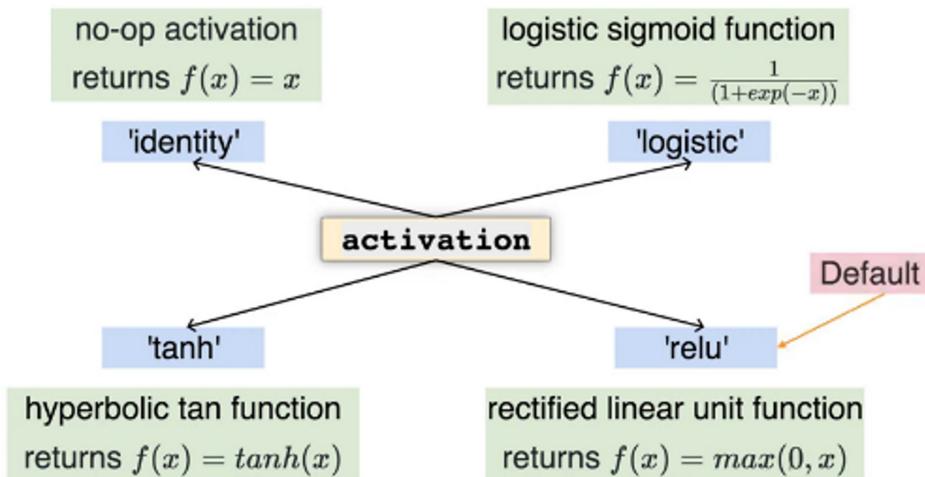


Default: 1 alpha = 0.0001 |

How to set the activation function for the hidden layers?

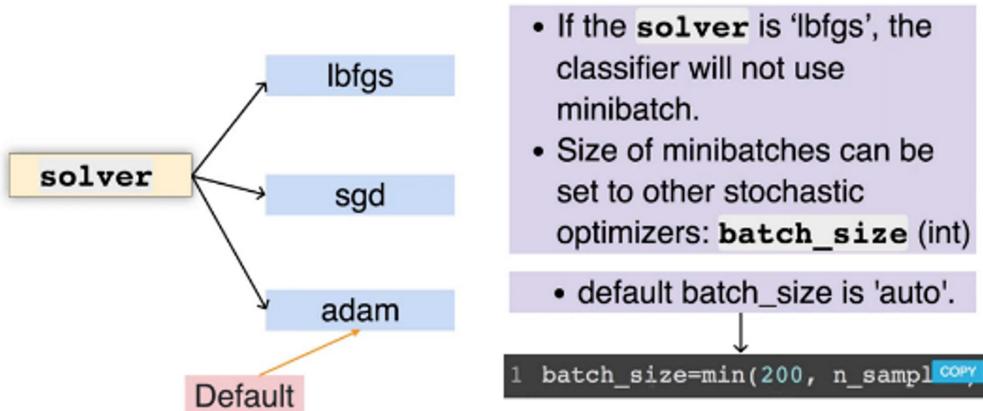


How to set the activation function for the hidden layers?



How to perform weight optimization in MLPClassifier?

- MLPClassifier optimizes the log-loss function using LBFGS or stochastic gradient descent



How to view weight matrix coefficients of trained MLPClassifier?

coefs_

- It is a **list** of shape (n_layers - 1,)
- The i th element in the list represents the weight matrix corresponding to layer i .

Example:

- "weights between input and first hidden layer:"

```
1 print(MLP_clf.coefs_[0])
```

```
weights between input and first hidden layer:  
[[-0.14201691 -1.18304359 -0.85567518 -4.53250719 -0.60466275]  
 [-0.69781111 -3.5850093 -0.26436018 -4.39161248 0.0644423]]
```

- "weights between first hidden and second hidden layer:"

```
1 print(MLP_clf.coefs_[1])
```

```
weights between first hidden and second hidden layer:  
[[ 0.29179638 -0.14155284]  
 [ 4.02666592 -0.61556475]  
 [-0.51677234 0.51479708]  
 [ 7.37215282 -0.31936965]  
 [ 0.32920668 0.64428109]]
```

How to view bias vector of trained MLPClassifier?



intercepts_

- It is a **list** of shape (n_layers - 1,)
- The i th element in the list bias vector corresponding to layer $i + 1$.

Example:

- "Bias values for first hidden layer:"

```
1 print(MLP_clf.intercepts_[0])
```

- "Bias values for second hidden layer:"

```
1 print(MLP_clf.intercepts_[1])
```

```
Bias values for first hidden layer:
```

Step 3: After fitting (training), the model can make predictions for new samples (X_{test}):

- returns predicted values for new samples
- for example:
`array([-0.9..., -7.1...])`

```
1 MLP_reg.predict(X_test)
```

- returns R^2 score

```
1 MLP_reg.score(X_test,y_test)
```

- for example:
`0.45678889`